

Sidhartha Thogari

Senior Data Engineer

@sid.siddhartha09@gmail.com

+44- 7774915173

United Kingdom

PROFESSIONAL SUMMARY

Databricks Certified Azure Data Engineer with 8+ years of experience designing, implementing, and maintaining enterprise-scale cloud data platforms, modern Lakehouse architectures, and analytics solutions. Strong expertise in Azure Databricks, PySpark, SQL, Azure Data Factory, Delta Lake, Unity Catalog, ETL/ELT, Data Warehousing, and Medallion Architecture. Skilled in leveraging modern Databricks capabilities including Delta Live Tables (DLT), Lakeflow, Databricks Workflows, AI/BI Dashboards, Databricks Genie, data lineage, data quality monitoring, and Lakehouse governance. Experienced in configuring Azure Data Factory (ADF), Azure Storage, Azure Event Hubs, Azure Key Vault, REST APIs, and Microsoft Graph API to integrate data from diverse enterprise systems. Proven experience in delivering scalable, secure, and high-performance data platforms while implementing data governance, automation, and GDPR-compliant solutions that enable trusted analytics and data-driven decision-making. Strong analytical, problem-solving, and communication skills with a proven ability to deliver reliable, business-focused data solutions.

KEY SKILLS

- **Expert in Enterprise Data Engineering, designing,** developing, and maintaining enterprise-scale cloud data platforms, modern Lakehouse architectures, scalable ETL/ELT pipelines, Medallion Architecture (Bronze, Silver, Gold), and enterprise data integration solutions using Azure Databricks, Azure Data Factory, PySpark, SQL, and Delta Lake.
- **Proficient in Modern Databricks Platform,** implementing Unity Catalog, Delta Live Tables (DLT), Lakeflow, Databricks Workflows, Databricks SQL, Auto Loader, AI/BI Dashboards, Databricks Genie, data lineage, data quality monitoring, and Lakehouse governance.
- **Experienced in Enterprise System Integration,** integrating data from ERP, CRM, eCommerce, POS, Digital Analytics, SharePoint, Microsoft Graph API, REST APIs, Azure Event Hubs, and other cloud-based enterprise applications.
- **Skilled in Customer Data Platforms and Data Governance,** designing and delivering Customer 360 and Single Customer View (SCV) solutions while implementing GDPR-compliant data governance, PII protection, Unity Catalog access controls, secure mapping tables, hashing frameworks, and Right-to-be-Forgotten (RTBF) processes.
- **Experienced in Cloud Automation and DevOps,** implementing Infrastructure as Code (IaC) using Terraform, Azure DevOps, Git, CI/CD pipelines, and automated deployment of secure and scalable enterprise cloud data platforms.
- **Strong in Performance Optimization and Collaboration,** optimising Spark workloads, SQL queries, Delta Lake performance, implementing automated monitoring, alerting, and data quality frameworks, while working closely with architects, analysts, and cross-functional teams to deliver reliable and business-ready data solutions.

TECHNICAL SKILLS

- **Languages:** Python, PySpark, R programming, C, C#, Java, CSS, HTML, XML
- **Databases:** MS SQL Server, MySQL, PostgreSQL, MongoDB, Cassandra
- **Cloud:** Microsoft Azure, Amazon AWS
- **Project Management tools:** MS Project, MS Visio, MS Office (Word, Excel, PowerPoint & Access), MS SharePoint
- **Methodologies:** Agile, SCRUM, Waterfall
- **Machine Learning Algorithms:** Decision Tree, Random Forest, K-NN, Naive Bayes, Logistic, Linear, NLP
- **Technologies:** Azure DevOps, Azure Databricks, Azure Data Factory (ADF), Azure Data Lake Storage

(ADLS Gen2), Azure Synapse Analytics, Terraform, ETL/ELT, Snowflake, Apache Spark, MLflow, Visual Studio Code, Apache Airflow, Matillion, Visual Studio, CI/CD Pipelines, Hadoop, Hive, Git, Data Warehousing, Data Modelling, Delta Lake, Delta Live Tables (DLT), Unity Catalog, Databricks Workflows, Lakeflow Pipelines, Auto Loader, Databricks SQL, AI/BI Dashboards, Databricks Genie, Data Lineage, Data Quality Monitoring, Medallion Architecture, Lakehouse Architecture, Performance Optimization, Data Pipeline Orchestration, Data Integration, SharePoint, Power Automate, Key Performance Indicators(KPIs)

PROFESSIONAL EXPERIENCE

Azure Data Engineer, Brompton Bicycle

Sept 2023 - Present

- Designed, developed, and maintained enterprise-scale Lakehouse platforms using Azure Databricks, PySpark, Delta Lake, and Medallion Architecture (Bronze, Silver, Gold), delivering scalable, secure, and high-performance data solutions.
- Built scalable batch and near real-time ETL/ELT pipelines using Delta Live Tables (DLT), Lakeflow, Auto Loader, Databricks Workflows, and Databricks SQL, automating data ingestion, transformation, and delivery across enterprise platforms.
- Implemented modern Databricks Lakehouse capabilities including Unity Catalog, AI/BI Dashboards, Databricks Genie, data lineage, data quality monitoring, and Lakehouse governance, enabling governed, AI-ready, and self-service analytics.
- Designed and implemented Customer 360 and Single Customer View (SCV) solutions by consolidating customer, sales, registration, subscription, and digital engagement data into trusted analytical datasets.
- Developed scalable enterprise data integration solutions by ingesting and transforming data from NetSuite, Salesforce, BigCommerce, Lightspeed, Smaregi, GA4, SharePoint, Microsoft Graph API, REST APIs, Oracle, Meetup, Jotform, Trade Platform, and other enterprise applications.
- Designed and implemented GDPR-compliant data governance frameworks including PII protection, hashing techniques, secure mapping tables, Unity Catalog access controls, and Right-to-be-Forgotten (RTBF) processes to ensure secure and compliant data management.
- Built AI-ready curated datasets, semantic data models, and trusted analytical data products to support Databricks Genie, AI/BI Dashboards, executive reporting, and self-service analytics across the organisation.
- Implemented automated data quality validation, monitoring, alerting, and operational support frameworks, improving production reliability, platform observability, and data integrity.
- Optimised Spark workloads, SQL queries, Delta Lake tables, and incremental processing strategies to improve pipeline performance, scalability, and operational efficiency across enterprise data platforms.
- Collaborated with business stakeholders, solution architects, analysts, and engineering teams to gather requirements, design scalable data architectures, and deliver business-ready data products that support strategic decision-making.

Data Engineer, Betway Group

Apr 2022 - Sept 2023

- Designed, developed, and maintained end-to-end ETL/ELT pipelines using Azure Data Factory (ADF), Azure Databricks, PySpark, and Delta Lake to ingest, transform, and load data from multiple enterprise data sources into Azure Data Lake.
- Developed scalable data processing solutions using Apache Spark, optimising ETL workloads and improving pipeline performance for large-scale analytical processing.
- Implemented Delta Lake and Delta Tables to provide ACID transactions, schema enforcement, version control, and high-performance data processing for enterprise data platforms.
- Built incremental data processing frameworks using Change Data Capture (CDC) techniques, reducing data latency and enabling efficient near real-time data synchronisation.
- Configured and optimised Azure Databricks clusters, Databricks Jobs, and automated workflow scheduling to improve processing efficiency, reliability, and cost optimisation.
- Developed real-time data ingestion pipelines using Azure Event Hubs and Azure Databricks to process streaming data from multiple business applications.

- Implemented Infrastructure as Code (IaC) using Terraform and developed automated CI/CD pipelines using Azure DevOps and YAML to streamline cloud infrastructure deployment and application delivery.

Data Engineer, Verizon

Feb 2020 - Apr 2022

- Developed and optimised end-to-end data pipelines, improving data processing efficiency by 30% through scalable data transformation and workflow optimisation.
- Designed and implemented data governance, security, and access control measures to ensure compliance with organisational data standards and privacy best practices.
- Optimised complex SQL queries using joins, indexing, partitioning, and query tuning techniques, improving query performance by 30% and reducing execution time.
- Developed SQL-based analytical solutions using aggregate functions, window functions, and data transformation techniques to generate business insights from large-scale datasets.
- Utilised Git for source code management, version control, and collaborative software development while following best coding practices.
- Worked in Agile (Scrum and Kanban) environments, collaborating with cross-functional teams to deliver high-quality data solutions within project timelines.

Associate Consultant, Capgemini

Sept 2018 – Feb 2020

- Designed and developed automation solutions using Python, Machine Learning, and Natural Language Processing (NLP), deploying applications on AWS EC2 and cloud environments.
- Developed AI-driven solutions for infrastructure monitoring, predictive server health analysis, and automated feedback systems to improve operational efficiency.
- Built Python-based data integration solutions, processing and loading data into AWS Cassandra, while deploying scalable applications using AWS Elastic Beanstalk.
- Developed data visualisations using Matplotlib, performed debugging, testing, and code optimisation to improve application reliability and performance.
- Collaborated with cross-functional teams to deliver application enhancements, streamline processes, and improve overall user experience using Agile methodologies.

Senior Software Engineer, Capgemini

Sept 2017 – Sept 2018

- Developed and enhanced enterprise web applications using C#, ASP.NET MVC, SQL Server, HTML, CSS, JavaScript, and AngularJS to deliver business requirements.
- Designed algorithms using R to improve application performance and reduce overall system costs.
- Resolved application defects, implemented enhancements, and improved system stability, performance, and user experience.
- Collaborated with development and testing teams throughout the software development lifecycle, contributing to application development, testing, deployment, and ongoing support.

CERTIFICATIONS

- [Databricks Certified Data Engineer Associate](#)
- [Azure Databricks Platform Architect](#)
- [Generative AI Fundamentals](#)
- [Databricks Lakehouse Fundamentals](#)
- Azure Data Engineer Associate (DP 203)
- Hashicorp Certified: Terraform Associate

EDUCATION

Master's in Applied Data Science
University of Essex, United Kingdom

Jan 2022